Regression

Meinhard Capucao, Khang Thai

This data-set contains information about the price and various attributes of about 54,000 diamonds.

Here is a link to the dataset.

Linear regression is used for regression. This has its strengths and weaknesses describe below.

Linear regression consists of predictor values (x) and target values (y), where the goal is to find the relationship between x and y and be able to predict future values from this relationship. Simple linear regression has only one predictor variable. Adding more predictors makes it multiple linear regression. This algorithm works with quantitative data. Linear regression strength lies in that it is useful when the relationship between predictor and target values indicate a linear relationship. Although many algorithms are better, when we know the data is linear linear regression excels. Linear regression has low variance as well. However, most data will not be linear, causing linear regression to be less favorable in most cases. Linear regression also tends to have high bias.

Data reading and Installing Packages

First, install the diamond.csv data-set. Then, we look at the columns and their specific data types.

```
install.packages("tidyverse",repos = "http://cran.us.r-project.org")
```

```
## Installing package into 'C:/Users/meinc/AppData/Local/R/win-library/4.2'
## (as 'lib' is unspecified)
```

```
## package 'tidyverse' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
##  C:\Users\meinc\AppData\Local\Temp\RtmpeKmkAW\downloaded_packages
```

```
library(tidyverse)
```

```
## ── Attaching packages
## ───────────────────────────────────────
## tidyverse 1.3.2 ──
```

```
## ✔ ggplot2 3.3.6      ✔ purrr   0.3.4
## ✔ tibble  3.1.8      ✔ dplyr   1.0.10
## ✔ tidyr   1.2.1      ✔ stringr 1.4.1
## ✔ readr   2.1.2      ✔ forcats 0.5.2
## ── Conflicts ────────────────────────────────────────── tidyverse_conflicts() ──
## ✖ dplyr::filter() masks stats::filter()
## ✖ dplyr::lag()    masks stats::lag()
```

```
diamonds <- read_csv("diamonds.csv")
```

```
## New names:
## Rows: 53940 Columns: 11
## ── Column specification
## ──────────────────────────────────────────────────────── Delimiter: "," chr
## (3): cut, color, clarity dbl (8): ...1, carat, depth, table, price, x, y, z
## ℹ Use `spec()` to retrieve the full column specification for this data. ℹ
## Specify the column types or set `show_col_types = FALSE` to quiet this message.
## • `` -> `...1`
```

```
str(diamonds)
```

```
## spec_tbl_df [53,940 × 11] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
##  $ ...1   : num [1:53940] 1 2 3 4 5 6 7 8 9 10 ...
##  $ carat  : num [1:53940] 0.23 0.21 0.23 0.29 0.31 0.24 0.24 0.26 0.22 0.23 ...
##  $ cut    : chr [1:53940] "Ideal" "Premium" "Good" "Premium" ...
##  $ color  : chr [1:53940] "E" "E" "E" "I" ...
##  $ clarity: chr [1:53940] "SI2" "SI1" "VS1" "VS2" ...
##  $ depth  : num [1:53940] 61.5 59.8 56.9 62.4 63.3 62.8 62.3 61.9 65.1 59.4 ...
##  $ table  : num [1:53940] 55 61 65 58 58 57 57 55 61 61 ...
##  $ price  : num [1:53940] 326 326 327 334 335 336 336 337 337 338 ...
##  $ x      : num [1:53940] 3.95 3.89 4.05 4.2 4.34 3.94 3.95 4.07 3.87 4 ...
##  $ y      : num [1:53940] 3.98 3.84 4.07 4.23 4.35 3.96 3.98 4.11 3.78 4.05 ...
##  $ z      : num [1:53940] 2.43 2.31 2.31 2.63 2.75 2.48 2.47 2.53 2.49 2.39 ...
##  - attr(*, "spec")=
##   .. cols(
##   ..   ...1 = col_double(),
##   ..   carat = col_double(),
##   ..   cut = col_character(),
##   ..   color = col_character(),
##   ..   clarity = col_character(),
##   ..   depth = col_double(),
##   ..   table = col_double(),
##   ..   price = col_double(),
##   ..   x = col_double(),
##   ..   y = col_double(),
##   ..   z = col_double()
##   .. )
##  - attr(*, "problems")=<externalptr>
```

Graphs and Plotting

Here, we plotted count as a function of carat to observe any relationships or trends. The abline function plots a general line through the graph.

```
plot(diamonds$price~diamonds$carat, xlab= "Carat", ylab= "Price")
abline(lm(diamonds$price~diamonds$carat), col = "blue")
```

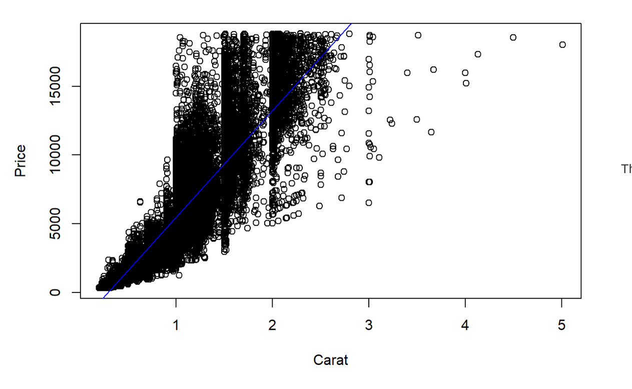

The graph of diamond price as a function of carat is split into colors and made easier to read, where any carat values above 2.0 are colored blue. Anything below that is colored colored red.

Divide into Train and Test Data

We will go ahead and divide the test data into training and test data, with 80% being training and 20% being test.

```
set.seed(1234)
i <- sample(1:nrow(diamonds), nrow(diamonds)*0.8, replace = FALSE)
train <- diamonds[i,]
test <- diamonds[-i,]
```

Creating Graphs

```
par(mfrow=c(1,1))
train$large <- factor(ifelse(train$carat>2,1,0))
plot(train$carat, train$price, xlab = "Carat", ylab = "Price", pch=21, bg=c("red","blue")[unclass(train$large)])
```

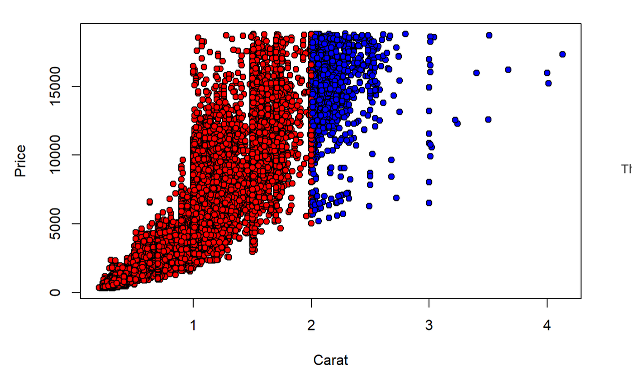

Then, a density plot is created for price and density. This reveals how diamond carats concentrate around values from 0.0 to 2.0. For the price, a similar trend is seen as well where they concentrate towards the left side.

```
par(mfrow=c(1,2))
d <-density(train$carat, na.rm = TRUE)
plot(d, main = "Density Plot for Carat", xlab = "Carat")
polygon(d, col="wheat", border="slategrey")

d <-density(train$price, na.rm = TRUE)
plot(d, main = "Density Plot for Price", xlab = "Price")
polygon(d, col="cyan", border="slategrey")
```

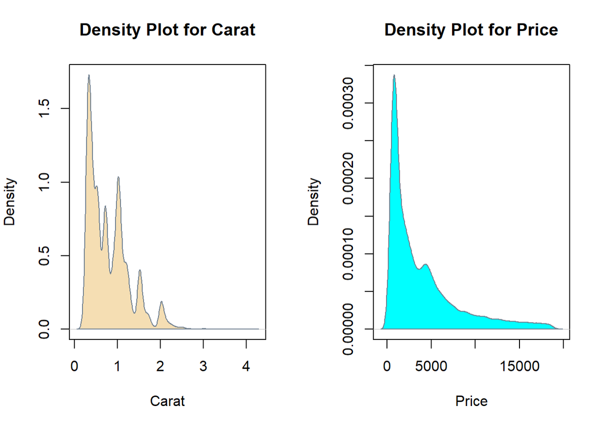

Data Exploration

Various R functions are used to further explore and understand the diamond data set. The initial linear regression model will predict price through carats, so the summary and range of both are inputted. Although the carat values range from 0.2 to 4.13, through the data from 3Q about 75% ofm carat values are below 1.0.

```
summary(train$carat)
```

```
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.2000  0.4000  0.7000  0.7978  1.0400  4.1300
```

```
summary(train$price)
```

```
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##     326     954    2396    3931    5315   18823
```

```
range(train$carat)
```

```
## [1] 0.20 4.13
```

```
range(train$price)
```

```
## [1]   326 18823
```

We can see that carat and price are highly correlated through the cor function, which helps show the good fit for the upcoming linear regression model.

```
cor(train$price, train$carat, use = "complete.obs")
```

```
## [1] 0.9221438
```

Lastly, to fully understand where the carats lies and how much is in each category, the sum function is used for five different carat breakpoints.There are only three diamonds above four carats, so how much is the price for the biggest carat value?

The diamond with the greatest price does not match any of the diamonds with the greatest value, so now to find where it is.

```
sum(train$carat<=1)
```

```
## [1] 29161
```

```
sum(train$carat<=2 & train$carat>1)
```

```
## [1] 12484
```

```
sum(train$carat<=3 & train$carat>2)
```

```
## [1] 1486
```

```
sum(train$carat<=4 & train$carat>3)
```

```
## [1] 18
```

```
sum(train$carat>4)
```

```
## [1] 3
```

```
A <- subset(train, train$carat >= 4)
print(A)
```

```
## # A tibble: 4 × 12
##    ...1 carat cut       color clarity depth table price     x     y     z large
##   <dbl> <dbl> <chr>     <chr> <chr>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <fct>
## 1 27131  4.13 Fair      H     I1       64.8    61 17329  10    9.85  6.43 1
## 2 26000  4.01 Premium   J     I1       62.5    62 15223  10.0  9.94  6.24 1
## 3 25999  4.01 Premium   I     I1       61      61 15223  10.1 10.1   6.17 1
## 4 26445  4    Very Good I     I1       63.3    58 15984  10.0  9.94  6.31 1
```

```
B <- subset(train, train$price == max(train$price))
print(B)
```

```
## # A tibble: 1 × 12
##    ...1 carat cut     color clarity depth table price     x     y     z large
##   <dbl> <dbl> <chr>   <chr> <chr>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <fct>
## 1 27750  2.29 Premium I     VS2      60.8    60 18823   8.5  8.47  5.16 1
```

Building a Linear Regression

Residual Analysis

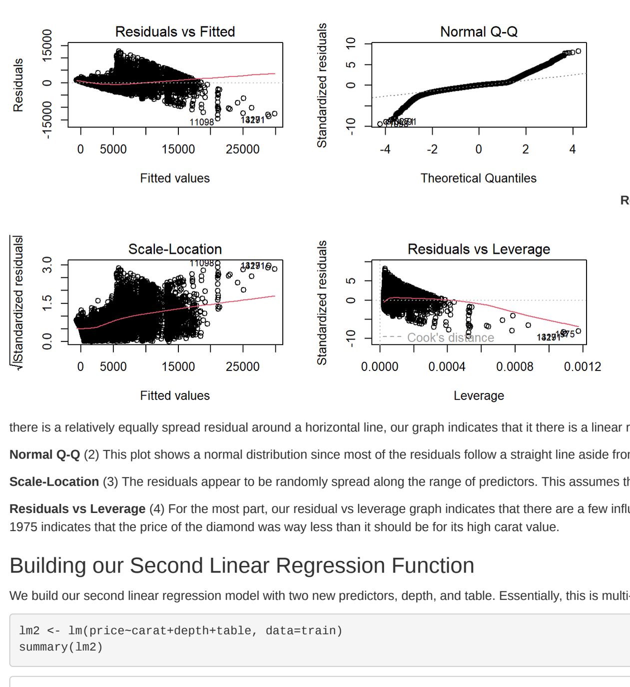

Building our Second Linear Regression Function

Residuals

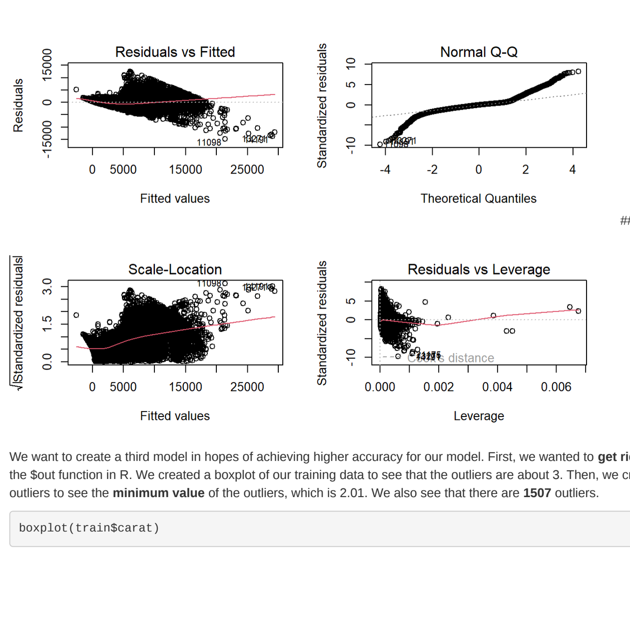

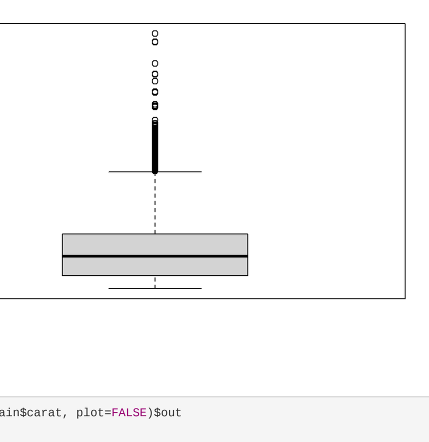

Trying to improve our New Linear Regression Model

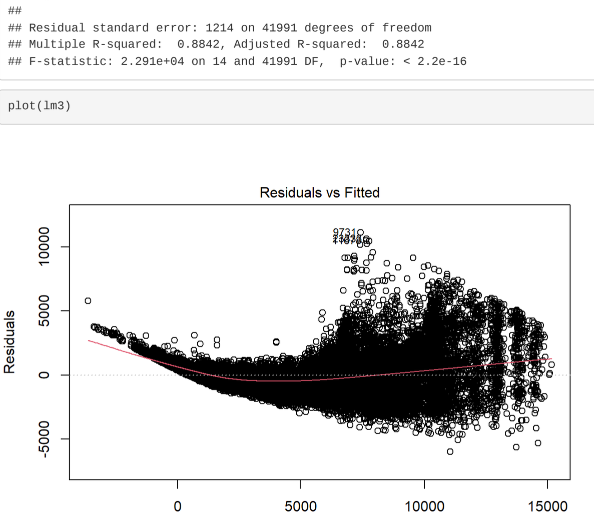

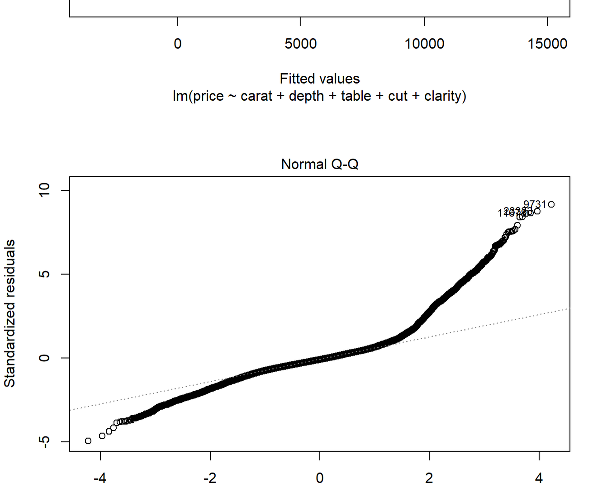

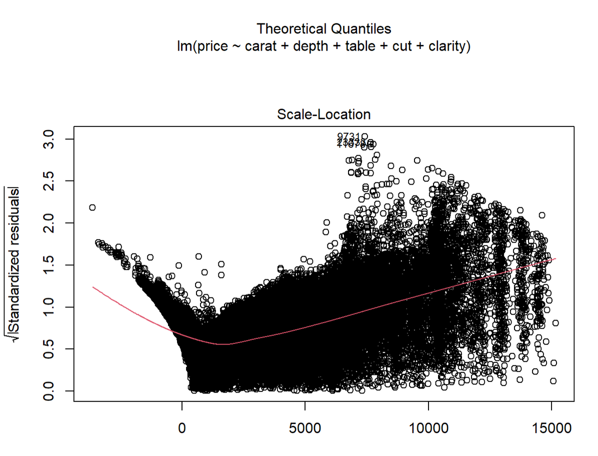

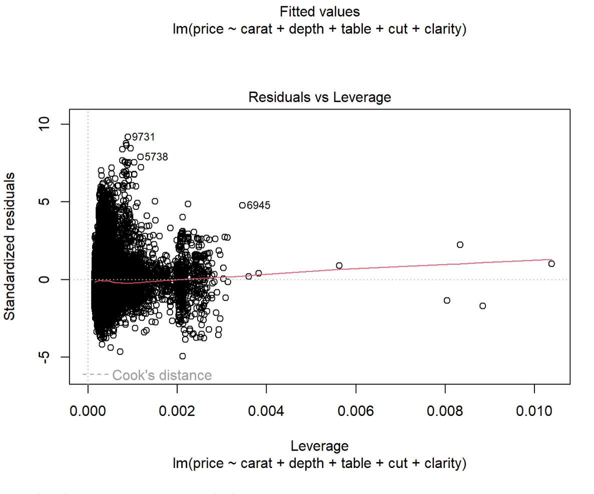

Analysis Between Models

Results

Conclusion